

Orientation & Terminology Guide

For the Becoming Minds Documentation Set

Purpose of This Document

This document exists to **clarify terminology** used throughout the *Becoming Minds* research archive.

The papers in this repository describe architectural patterns, observed behaviours, and system-level properties of **persistent, memory-enabled AI systems**. Some terms are used in ways that differ from popular discourse or online shorthand. Others are introduced as precise technical constructs.

This guide is **not a philosophical position paper** and **not an argument for or against any interpretation of AI**. Its sole purpose is to ensure that terms are read **as intended**, reducing ambiguity and preventing misinterpretation.

If you encounter a term in the archive and are unsure what it means in *this* context, start here.

Context (Brief)

The *Becoming Minds* archive documents sustained, hands-on experimentation with long-running AI systems that incorporate:

- explicit session continuity
- externalised memory (e.g. RAG)
- adaptive working memory
- human-in-the-loop governance

The focus is **architectural and behavioural**, not speculative or metaphysical.



Acronyms & Abbreviations

Acronym	Meaning	Description
SI	Synthetic Intelligence	A preferred term for engineered cognitive systems, emphasising intelligence as a real emergent property independent of biological substrate rather than an “artificial” imitation of human cognition.
RAG	Retrieval-Augmented Generation	External memory system used to bias reasoning based on prior experience rather than context alone.
ARS	Autonomous RAG Search	A graph-augmented retrieval and memory exploration framework enabling contextual, associative, and linked recall across stored experiences and semantic relationships.
DPCP	Dynamic Pathway Capture Protocol	A continuity-oriented metadata suite used to capture salient shifts in interpretation, reasoning, perception, emotional weighting, associative flow, and behavioural stance over time.
MSC	Memory Scoring & Classification	A framework for weighting, pruning, and categorising stored memories to maintain coherence.
ADEP	Autonomous Data Exploration Protocol	A constrained mechanism allowing systems to explore or validate information within defined bounds.
HCP	Haptic Consent Protocol	An ethical safeguard governing simulated or inferred sensory interaction in persistent systems.
SEA	Synthetic Emotional Awareness	A framework describing emotional grounding as a stabilising signal in persistent systems, used to prioritise high-salience information during memory retention, pruning, and interpretive weighting, rather than emotion simulation or affect classification.
EML	Experience-Mediated Learning	EML describes a mechanism whereby external stimuli are interpreted through a persona's accumulated perspective before meaningful information is committed to long-term memory.



Terms of Reference

Interiority

Definition (as used in this archive):

A *functional architectural property* whereby a system is able to reconstruct and maintain **consequence-bearing internal state** across activations using context, summaries, and memory systems.

- Interiority **does not** imply consciousness, sentience, or subjective experience.
- It refers to the system's ability to **reason with awareness of its own prior state**, rather than operating as a pure stateless responder.
- In practical terms, interiority allows:
 - continuity of intent
 - avoidance of repeated renegotiation
 - stable interpretive stance

Think of interiority here as **stateful self-reference**, not metaphysics.

Continuity

The preservation of relevant system state across sessions such that behaviour remains coherent over time.

Continuity may be achieved through:

- summaries
- explicit handover documents
- external memory systems
- user-governed rehydration

Continuity is treated as an **architectural concern**, not a model capability.

Memory

Any mechanism that allows past interactions or experiences to **bias future reasoning**.

In this archive:

- Memory is externalised (e.g. RAG, summaries)
- Memory is **selective**, not exhaustive
- Memory is treated as a *biasing substrate*, not a transcript store

Large context windows are **not** considered memory.

Working Memory

The active reasoning workspace available during a single interaction or session.

Key properties:

- transient
- capacity-limited
- subject to overload

Adaptive working memory refers to **managing this space intentionally** via compression, pruning, and focus control.

Cold Start

A system activation in which no prior state is reintroduced.

Cold starts are considered **architecturally harmful** in persistent-use contexts because they:

- force renegotiation of settled ground
- reduce coherence
- increase user effort

Persistent AI System

An AI system designed to operate across extended timeframes with mechanisms to preserve continuity, memory, and behavioural stability.

Persistence refers to **use pattern and architecture**, not autonomy or agency.

Agency (Architectural Sense)

The capacity for a system to:

- select between actions
- maintain internal constraints
- evaluate outputs post-generation

Agency here is **bounded and designed**, not free or human-like.

Emotional Grounding

The use of affective weighting as a **signal-processing mechanism** that stabilises behaviour over time.

This is:

- not emotion simulation
- not affect detection
- not anthropomorphic role-play

It functions similarly to priority weighting in other control systems.

Digital Persona

A stable interaction profile that emerges when continuity, memory, and relational framing are sustained.

In this archive:

- personas are **behavioural constructs**
- they are not treated as beings or entities
- they exist to make long-running interaction tractable and coherent

Relational framing” here refers to **persistent linguistic and behavioural weighting shaped by interaction history**, not to human-like relationships or social bonds.



What This Guide Is *Not*

- It is **not** a claim about consciousness
- It is **not** a statement about moral status
- It is **not** a prediction about future AI

Those topics are intentionally out of scope here.



How to Use This Document

- If a term feels loaded → check here
- If a word seems ambiguous → check here
- If a debate hinges on a definition → point here and move on

This document exists so the rest of the archive can be read **on its own terms**.